

CS 598 PS

Machine Learning for Signal Processing

Problem Set 2

Christian Howard
howard28@illinois.edu

Abstract

Within this report is a set of solutions for Problem Set 2 spanning areas of linear and nonlinear unsupervised learning applied to sound and images. In the first problem, Principle Component Analysis (PCA), Independent Component Analysis (ICA), and Non-negative Matrix Factorization (NMF) were applied to identifying three different types of sounds played in a recording based on a spectrogram of the sound signal. The second problem used PCA, ICA, and NMF to again identify features, this time for a dataset of hand written digits. In the first two problems, it was found that NMF produced more intuitive features and weights that were purely positive and in turn additive. Additionally, ICA tended to create a much more independent looking set of features than the PCA data it was based on, as expected. In the third problem, PCA and Laplacian Eigenmaps were used to try and understand the structure of the digit 6 subset. It was found that PCA did not do a great job of showing the nonlinear structure of this data while Laplacian Eigenmaps did very well at achieving this.

Contents

1	Problem 1 - An Audio Features Project	3
2	Problem 2 - Handwritten Digit Features	13
3	Problem 3 - The Geometry of Handwritten Digits	15
4	Software Used	18
5	Specialized Software Written	19

1 Problem 1 - An Audio Features Project

Within this problem, the goal is to use a set of dimensionality reduction and feature selection techniques to find three features that can help identify three different sounds in a recording. The methods being reviewed are PCA, ICA, Non-negative Matrix Factorization (NMF), and then for fun I wrote an algorithm, call is the Randomized Projection (RP), that forms a low-rank approximation via projection form to the data D using randomized algorithms such that $D = Q(Q^T D) = QB$ and $Q^T Q = I$. Before jumping into the analysis with the various feature selection methods, Figure 1 represents a visual of the spectrogram for the input sound which will be the data set being passed into the set of algorithms.




Figure 1: Spectrogram of Sound

With Figure 1 representing the input data set, one can then produce the

three features using each method. Figures 2 and 3 represent the features and weights generated by PCA. Figures 6 and 7 represent the features and weights generated by chaining PCA and ICA. Figures 8 and 9 represent the features and weights generated by NMF. Figures 4 and 5 represent the features and weights generated by the RP approach. In terms of observations, all of the methods produce features with generally similar patterns. However, ICA definitely produces features that appear more independent than those of PCA. The RP approach also produces a very similar set of features to PCA but without having to compute an expensive SVD.

Additionally, PCA and ICA have features that have components that are negative and positive, while NMF has features that have purely positive components. This difference produces interesting behavior in the weights because PCA and ICA have both negative and non-negative weights, while NMF produces purely positive weights. At least for myself, NMF produces the pair of features and weights that make the most sense. This is not only because you can view the sound frequencies as purely additive, which is more intuitive, but when the weights are most positive, you can easily tell that is when a feature is most activated. For ICA and PCA, it is harder due to the mix of negative and positive weights and feature components.




Figure 2: PCA Features




Figure 3: PCA Weights




Figure 4: Randomized Projection Features




Figure 5: Randomized Projection Weights




Figure 6: ICA Features




Figure 7: ICA Weights




Figure 8: NMF Features




Figure 9: NMF Weights

2 Problem 2 - Handwritten Digit Features

Within this section, we look at finding features for a data set of hand drawn digits. In particular, we are tasked to find 36 features using PCA, ICA, and NMF. Figures 10, 11, and 12 display features found using PCA, ICA, and NMF.

As in Problem 1, PCA and ICA produce features that share the trait of having components that range from being positive to negative while NMF has features with purely positive components. This difference makes NMF produce features that are very different from PCA and ICA, again allowing NMF's features to be purely additive and fairly intuitive. When comparing PCA and ICA, we can see that ICA has features that appear more independent than that of PCA. This is because ICA's features tend to be more distinct features while PCA has features that share similarities. This is obviously expected since ICA strives to obtain independent features while PCA is just looking to create decorrelated features.




Figure 10: PCA Digit Features




Figure 11: ICA Digit Features




Figure 12: NMF Digit Features

3 Problem 3 - The Geometry of Handwritten Digits

Within this problem, we extend work done in Problem 2 by investigating patterns for the digit 6 in the data set. We start off this investigation by performing PCA on this subset of data to find 2 features, shown below in Figure 13.




Figure 13: PCA Digit 6 Features

If we take weights associated with the above features, we can plot them in a 2D scatter plot with their associated image from the original data set. We can see the result of such a thing in Figure 14.

As we look at Figure 14, we can see there appears to be a pattern going from left to right in the similarity between images. To see if there is some structure we can identify, we can use a Spectral Embedding via a Laplacian Eigenmap to see if we can find some underlying manifold for the data. For this case, we will produce this Laplacian Eigenmap using 10 nearest neighbors and again ensure we end up with two dimensional features. Figure 15 shows the resulting 2D scatter plot of the data set's images tied to their associated weight vectors.

As we can readily see as we compare Figure 14 to Figure 15, there indeed appears to be some underlying nonlinear structure for the 6 digit data set that has become more apparent after the embedding. This shows there is indeed a strong, nonlinear relationship between the data for the digit 6 that has been found via the manifold the Laplacian Eigenmap estimated.




Figure 14: PCA Weight Space Scatter Plot




Figure 15: Laplacian Eigenmap Weight Space Scatter Plot

4 Software Used

The following software was used to help in computing the various factorizations:

- scikit-learn for ICA, Non-negative Matrix Factorization, and the Laplacian Eigenmap computations

5 Specialized Software Written

```
1 # import important libs
2 import numpy as np
3 import numpy.linalg as la
4
5
6 def pca( D, k_or_tol ):
7     # Author: C. Howard
8     # method to perform PCA on some input data given the dimensionality
9     # we want our resulting features to have or given a tolerance on the
10    # ratio between singular values and the max singular value
11
12    # perform SVD  $\ni D = U\Sigma V^T$ 
13    [U,S,Vt] = la.svd(D, full_matrices = False )
14
15    # make copy of singular values
16    iv = np.copy(S)
17
18    # define current number of weight terms as follow
19    k = k_or_tol
20
21    # if we are picking k using a relative difference
22    # between singular value magnitudes
23    if k_or_tol < 1:
24        b = iv >= S[0]*k_or_tol # find indices  $j = 1, \dots, k \ni \sigma_j > \epsilon$ 
25        k = np.sum(b)
26
27    # compute  $W$  and  $W^+$  using truncated SVD terms
28    W      = U[:, :k].T
29    Winv   = U[:, :k]
30
31    # return weights and original singular values
32    return (W,Winv,S)
```

```

1 #import useful libraries
2 import numpy as np
3
4 def ica( D , eps = 1e-6, random_seed = 17):
5     # Author: C. Howard
6     # method to perform ICA using FastICA on some input data to create a set
7     # of independent features
8     #
9     # Inputs:
10    # D : Dataset where columns represent samples and
11    # rows are number of dimensions
12    # eps : the epsilon used to denote when the algorithm has converged
13    # random_seed : seed to use to make algorithm repeatable
14    #
15    # Outputs:
16    # W : Mixing Matrix
17    # Winv: Inverse of Mixing Matrix
18    #
19    # Code based on https://en.wikipedia.org/wiki/FastICA and comments made in
20    # Machine Learning a Probabilistic Perspective by Kevin P. Murphy
21
22    # initialize random seed for repeatability
23    np.random.seed(seed=random_seed)
24
25    # define the g(u) and g'(u) that will be used
26    def g_deriv(u):
27        c = np.tanh(u)
28        return (c, 1 - c**2)
29
30    # get number of components that will be used
31    (N,M) = D.shape
32    ncomp = N
33
34    # define output matrix W
35    W = np.random.randn(D.shape[0],ncomp)
36
37    # define some helper entities
38    One = np.ones((M,1))
39
40    # compute the components
41    for p in range(0,ncomp):
42
43        # set difference between w vectors to arbitrary number above threshold
44        dw = 10*eps
45

```

```

46     while dw > eps:
47
48         # get current value for vector w
49         w0 = np.copy(W[:,p])
50
51         # compute source signal
52         z = np.matmul(W[:,p].T,D)
53
54         # compute g and g' given z
55         (g,gd) = g_gderiv(z)
56
57         # set the new value for the vector w
58         W[:,p] = (np.matmul(D,g.T) - np.matmul(gd,One)*W[:,p])/M
59
60         # orthogonalize relative to the past w components
61         if p > 0:
62             W[:,p] = W[:,p] - np.matmul( W[:,0:p], (np.matmul(W[:,0:p].T,W[:,p])) )
63
64         # normalize to keep on constraint surface
65         w1      = W[:,p]
66         W[:,p]  = W[:,p]/np.sqrt(np.matmul(W[:,p].T,W[:,p]))
67         w2      = W[:,p]
68
69         # compute change in w vector
70         dw = np.linalg.norm(w0-W[:,p])
71
72     # Compute inverse of W
73     Winv = W.T
74
75     # return W and inverse of W
76     return (W,Winv)

```

```

1 #import useful libraries
2 import numpy as np
3
4 def projrep(A, k_or_tol, random_seed = 17, num_power_method = 4):
5     # Author: Christian Howard
6     # This function is designed to take some input matrix A
7     # and approximate it by the Low-rank form  $A = Q*(Q^T*A) = Q*B$ .
8     # This form is achieved using randomized algorithms and
9     # allows for adaptive rank reduction
10    #
11    # Inputs:
12    # A: Matrix to be approximated by Low-rank form
13    # k_or_tol: If value >= 1, it sets that rank as the target.
14    #           If 0 < value < 1, tries to adaptively find low rank form
15
16    # get dimensions of A
17    (r, c) = A.shape
18
19    # get the smallest dimension
20    sdim = min(r, c)
21
22    if k_or_tol >= 1:
23        k = int(k_or_tol)
24
25        # get the random input and measurements from column space
26        omega = np.random.randn(c, k)
27        Y = np.matmul(A, omega)
28
29        # form estimate for Q using power method
30        for i in range(1, num_power_method):
31            Q1, R1 = np.linalg.qr(Y)
32            Q2, R2 = np.linalg.qr(np.matmul(A.T, Q1[:, :k]))
33            Y = np.matmul(A, Q2[:, :k])
34            Q3, R3 = np.linalg.qr(Y)
35
36        # get final k orthogonal vector estimates from column space
37        Q = Q3[:, :k]
38
39        # compute weights associated with the column space to approximate A
40        B = np.matmul(Q.T, A)
41
42        # return the two matrices
43        return (Q, B)
44    else:
45

```

```

46     # init some variables used in algorithm
47     eps = k_or_tol
48     err = 1e20
49     kt = 0
50     k = math.ceil(0.14 * sdim)
51     iter_max = 14
52     iter = 1
53     Qt = np.zeros((r, c))
54
55     # adaptively try to estimate the rank via
56     # finding column space estimates
57     while err > eps and iter_max != iter:
58
59         # get the random input and measurements
60         omega = np.random.randn(c, k)
61         Y = np.matmul(A, omega)
62
63         # form estimate for Q using power method
64         for i in range(1, num_power_method):
65             Q1, R1 = np.linalg.qr(Y)
66             Q2, R2 = np.linalg.qr(np.matmul(A.T, Q1[:, :k]))
67             Y = np.matmul(A, Q2[:, :k])
68             Q3, R3 = np.linalg.qr(Y)
69             Q = Q3[:, :k] # get final k orthogonal vector estimates from column space
70
71         # compute error estimate
72         o = Q[:, 0] # use eigenvector from A as initial "random" vec
73
74         # compute normalized eigenvector estimate for error matrix E
75         z = np.matmul(A, o)
76         y = z - np.matmul(Q, np.matmul(Q.T, z))
77         y = y / np.amax(y)
78
79         # use multiple iterations of power method  $y = (E^*E')^{\{n\}}*y$ 
80         # to get most dominant singular vector
81         for i in range(1, num_power_method):
82             z = np.matmul(A.T, y)
83             y = z - np.matmul(Q, np.matmul(Q.T, z))
84             y = y / np.amax(y)
85             z = np.matmul(A, y)
86             y = z - np.matmul(Q, np.matmul(Q.T, z))
87             y = y / np.amax(y)
88
89         # compute largest singular value estimate based on dominant singular vector
90         # assume singular value estimate as the error since same as 2-norm of E
91         z = np.matmul(A, y)

```

```

92     err = abs(np.dot(y, z - np.matmul(Q, np.matmul(Q.T, z))) / np.dot(y, y))
93
94     # Stitch new Q vectors to net Q vectors
95     Qt[:, kt:(kt + k)] = Q
96
97     # update total k value
98     kt = kt + k
99
100    # orthogonalize vectors using QR
101    if iter != 1:
102        Qh, Rh = np.linalg.qr(Qt[:, :kt])
103        Qt[:, :kt] = Qh
104
105    # retrieve modified k-recent Q vectors
106    Q = Qt[:, (kt - k):kt]
107
108    # compute next best value for k
109    # current strategy uses minimized regret method based on solution to
110    # the egg drop algorithm problem
111    k = min(sdim, k + math.ceil((iter_max - iter) * sdim / 100.0)) - k
112    iter += 1
113
114    # update matrix A based on new Q vectors
115    # Note that new A matrices are essentially error matrices
116    # since we subtract approximate projection representations of the kth A matrix
117    A = A - np.matmul(Q, np.matmul(Q.T, A))
118
119    # set the final Q matrix
120    Q = Qt[:, :kt]
121
122    # compute weights associated with the column space to approximate A
123    B = np.matmul(Q.T, A)
124
125    # return the two matrices
126    return (Q, B)

```

```

1  #import useful libraries
2  import numpy as np
3  import math
4
5  def spectrogram(signal, ws = 1024, hs = 512):
6      # Author: C. Howard
7      # Function to compute the baseline complex valued spectrogram for some sound.
8      # signal: a sound represented as a column vector
9      # ws      : the window size
10     # hs      : the hope size, aka the amount we shift from sample to sample
11
12     # compute Hamming weights
13     alpha = 0.54
14     beta  = 1 - alpha
15     pi2   = 2.0*math.pi
16     c     = pi2/(ws-1)
17     w     = alpha - beta * np.cos(c*np.arange(0,ws,1))
18
19     # compute DFT matrix
20     p = np.arange(0,ws,1).reshape(ws,1)
21     F = np.exp( -(pi2/ws)*np.matmul(p,p.T)*1j ) / np.sqrt(ws)
22
23     # compute resulting local matrix
24     D = np.multiply(F,w)
25
26     # Compute number of samples in spectrogram
27     (len,c) = signal.shape
28     num_samples = math.floor((len - hs)/(ws - hs))
29
30     # initialize output S
31     S = np.zeros((ws,num_samples),dtype=type(F[0,0]))
32
33     # Loop through and construct S
34     for i in range(0,num_samples):
35         S[:ws, i] = np.matmul(D,signal[i*hs:(i*hs+ws),0])
36
37     # return the output spectrogram
38     return S[:int(ws/2),:]

```
